

Functions HW

```
# Initial packages required (we'll be adding more)
library(tidyverse)
library(nycflights13)
```

Filter/remove the NA values before rescaling

```
df <- tibble(
  a = rnorm(10),
  b = rnorm(10),
  c = rnorm(10),
  d = rnorm(10)
)
df$a[5] = NA
df
```

```
# A tibble: 10 x 4
      a      b      c      d
  <dbl> <dbl> <dbl> <dbl>
1 -1.10 -1.06  0.442 -0.115
2 -1.04  1.54 -0.237 -1.52
3 -0.690 -0.649 -2.50  0.0529
4 -0.634  0.249  1.87 -0.848
5 NA     -1.16  1.68  0.156
6  0.775 -0.427 -0.0000380 -0.00459
7  2.10 -0.0398 -0.418 -0.0274
8 -1.88  1.20  0.359 -0.329
9 -0.635  0.808 -1.48 -0.00898
10  0.815 -0.815  0.0434 -0.607
```

```
rescale_basic <- function(x) {
  (x - min(x)) / (max(x) - min(x))
}
```

```
df |>
  filter(!is.na(a)) |>
  mutate(new_a = rescale_basic(a))
```

```
# A tibble: 9 x 5
      a      b      c      d new_a
  <dbl> <dbl> <dbl> <dbl> <dbl>
1 -1.10 -1.06  0.442 -0.115  0.198
2 -1.04  1.54 -0.237 -1.52   0.212
3 -0.690 -0.649 -2.50   0.0529  0.300
4 -0.634  0.249  1.87   -0.848  0.314
5  0.775 -0.427 -0.0000380 -0.00459 0.667
6  2.10 -0.0398 -0.418   -0.0274  1
7 -1.88  1.20  0.359   -0.329  0
8 -0.635  0.808 -1.48   -0.00898 0.314
9  0.815 -0.815  0.0434  -0.607  0.678
```

[Pause to Ponder:] Do you notice anything in the output above that gives you pause?

The filter removed the entire row for the NA value.

Create an if statement to check if there are NAs; return an error if NAs exist

First, here's an example involving weighted means:

```
# Create function to calculate weighted mean
wt_mean <- function(x, w) {
  sum(x * w) / sum(w)
}
wt_mean(c(1, 10), c(1/3, 2/3))
```

```
[1] 7
```

```
# ((1 * 1/3) + (10 * 2/3)) / (1/3 + 2/3)

wt_mean(1:6, 1:3)
```

```
[1] 7.666667
```

```
# ((1 * 1) + (2 * 2) + (3 * 3) + (4 * 1) + (5 * 2) + (6 * 3)) / (1 + 2 + 3)
```

[Pause to Ponder:] Why is the answer to the last call above 7.67? Aren't we taking a weighted mean of 1-6, all of which are below 7?

The `w` input (1:3) is being recycled in the numerator twice, but is not being recycled at all in the denominator. Shows you can have code with no problems but will give you the wrong answer with no alerts.

```
# update function to handle cases where data and weights of unequal length
wt_mean <- function(x, w) {
  if (length(x) != length(w)) {
    stop("`x` and `w` must be the same length", call. = FALSE)
  } else {
    sum(w * x) / sum(w)
  }
}
wt_mean(1:6, 1:3)
```

Error:

```
! `x` and `w` must be the same length
```

```
# should produce an error now if weights and data different lengths
# - nice example of if and else
```

[Pause to Ponder:] What does the `call.` option do?

If true, it will pinpoint exactly which function produced the error

Now let's apply this to our rescaling function

```
rescale_w_error <- function(x) {
  if (is.na(sum(x))) {
    stop("`x` cannot have NAs", call. = FALSE)
  } else {
    (x - min(x)) / (max(x) - min(x))
  }
}
```

```
temp <- c(4, 6, 8, 9)
rescale_w_error(temp)
```

```
[1] 0.0 0.4 0.8 1.0
```

```
temp <- c(4, 6, 8, 9, NA)
rescale_w_error(temp)
```

```
Error:
! `x` cannot have NAs
```

[Pause to Ponder:] Why can't we just use `if (is.na(x))` instead of `is.na(sum(x))`?

We have to use the latter because it prints a number of the total NA values while `is.na(x)` prints a logical vector which has a length longer than 1. The condition cannot have a length greater than 1 in order for the function to run.

[Pause to Ponder:] Predict what the code below will do, and (only) then run it to check. Think about: why do we have `sort = sort`? why not embrace `df`? why didn't we need `n` in the arguments?

The code below will make a function that will filter a data frame by a condition and then count the levels of the chosen variable in order to calculate the proportion of each level. We need `sort = sort` because the second `sort` will be the actual argument stated (TRUE or FALSE). The `sort` is there for the final output of proportions: if the user of the function does not want them sorted from higher to lower, they can state that by coding `sort = FALSE`. We don't need to embrace `df` because it is not within a function that requires embracing. We didn't need `n` in the arguments because it was produced by the count function within the new function.

```
new_function <- function(df, var, condition, sort = TRUE) {
  df |>
    filter({{ condition }}) |>
    count({{ var }}, sort = sort) |>
    mutate(prop = n / sum(n))
}

mpg |> new_function(var = manufacturer,
                    condition = manufacturer %in% c("audi",
                                                    "honda"),
```

```

)
"hyundai",
"nissan",
"subaru",
"toyota",
"volkswagen")

```

[Pause to Ponder:] Here's another nice use of `pick()`. Predict what the function will do, then run the code to see if you are correct.

The new function will count both the rows and columns of the data and pivot the data to have all the values from the column to be new columns with their new values being the counts from above.

After running: I do not know what happened honestly.

```

# Source: https://twitter.com/pollicipes/status/1571606508944719876
new_function <- function(data, rows, cols) {
  data |>
    count(pick(c({{ rows }}), {{ cols }}))) |>
    pivot_wider(
      names_from = {{ cols }},
      values_from = n,
      names_sort = TRUE,
      values_fill = 0
    )
}

mpg |> new_function(c(manufacturer, model), cyl)

# figuring it out?
mpg |>
  count(pick(c(c(manufacturer, model)), cyl)) # counts manufacturer by model and cyl?

mpg |>
  count(pick(c(c(manufacturer, model)), cyl)) |>
  pivot_wider(
    names_from = cyl, # new column names
    values_from = n, # values from the number cars (?) from each model and manufacturer
    names_sort = TRUE,
    values_fill = 0
  )

```

On Your Own

1. Rewrite this code snippet as a function: `x / sum(x, na.rm = TRUE)`. This code creates weights which sum to 1, where NA values are ignored. Test it for at least two different vectors. (Make sure at least one has NAs!)

```
weight01 <- function(x, na.rm = FALSE) {  
  x / sum(x, na.rm = na.rm)  
}
```

```
weight01(c(1:5, NA), na.rm = TRUE)
```

```
[1] 0.06666667 0.13333333 0.20000000 0.26666667 0.33333333      NA
```

```
weight01(10:20, na.rm = TRUE)
```

```
[1] 0.06060606 0.06666667 0.07272727 0.07878788 0.08484848 0.09090909  
[7] 0.09696970 0.10303030 0.10909091 0.11515152 0.12121212
```

2. Create a function to calculate the standard error of a variable, where $SE = \text{square root of the variance divided by the sample size}$. Hint: start with a vector like `x <- 0:5` or `x <- gss_cat$age` and write code to find the SE of `x`, then turn it into a function to handle any vector `x`. Note: `var` is the function to find variance in R and `sqrt` does square root. `length` may also be handy. Test your function on two vectors that do not include NAs (i.e. do **not** worry about removing NAs at this point).

```
se <- function(x) {  
  sqrt(var(x) / length(x))  
}
```

```
se(1:5)
```

```
[1] 0.7071068
```

3. Use your `se` function within `summarize` to get a table of the mean and s.e. of `hwy` and `cty` by `class` in the `mpg` dataset.

```
mpg |>
  group_by(class) |>
  summarize(
    mean_hwy = mean(hwy),
    se_hwy = se(hwy),
    mean_cty = mean(cty),
    se_cty = se(cty)
  )
```

```
# A tibble: 7 x 5
  class      mean_hwy se_hwy mean_cty se_cty
<chr>      <dbl>   <dbl>   <dbl>   <dbl>
1 2seater    24.8   0.583    15.4   0.245
2 compact    28.3   0.552    20.1   0.494
3 midsize    27.3   0.334    18.8   0.304
4 minivan    22.4   0.622    15.8   0.553
5 pickup     16.9   0.396     13    0.356
6 subcompact  28.1   0.909    20.4   0.778
7 suv        18.1   0.378    13.5   0.307
```

4. Use your `se` function within `summarize` to get a table of the mean and s.e. of `arr_delay` and `dep_delay` by carrier in the `flights` dataset. Why does the output look like this?

Because NAs are included in the flights data so mean and se can't be calculated.

```
flights |>
  group_by(carrier) |>
  summarize(
    mean_arr_delay = mean(arr_delay),
    se_arr_delay = se(arr_delay),
    mean_dep_delay = mean(dep_delay),
    se_dep_delay = se(dep_delay)
  )
```

```
# A tibble: 16 x 5
  carrier mean_arr_delay se_arr_delay mean_dep_delay se_dep_delay
<chr>      <dbl>      <dbl>      <dbl>      <dbl>
1 9E          NA          NA          NA          NA
2 AA          NA          NA          NA          NA
3 AS          NA          NA          NA          NA
4 B6          NA          NA          NA          NA
5 DL          NA          NA          NA          NA
```

6 EV	NA	NA	NA	NA
7 F9	NA	NA	NA	NA
8 FL	NA	NA	NA	NA
9 HA	-6.92	4.06	4.90	4.01
10 MQ	NA	NA	NA	NA
11 OO	NA	NA	NA	NA
12 UA	NA	NA	NA	NA
13 US	NA	NA	NA	NA
14 VX	NA	NA	NA	NA
15 WN	NA	NA	NA	NA
16 YV	NA	NA	NA	NA

5. Make your `se` function handle NAs with an `na.rm` option. Test your new function (you can call it `se` again) on a vector that doesn't include NA and on the same vector with an added NA. **Be sure to check that it gives the expected output with `na.rm = TRUE` and `na.rm = FALSE`.** Make `na.rm = FALSE` the default value. Repeat #4. (Hint: be sure when you divide by sample size you don't count any NAs)

```
se <- function(x, includeNAs) {
  if (includeNAs) {
    x <- na.omit(x)
  }
  sqrt(var(x) / length(x))
}

flights |>
  group_by(carrier) |>
  summarize(
    mean_arr_delay = mean(arr_delay, na.rm = TRUE),
    se_arr_delay = se(arr_delay, includeNAs = TRUE),
    mean_dep_delay = mean(dep_delay, na.rm = TRUE),
    se_dep_delay = se(dep_delay, includeNAs = TRUE)
  )
```

A tibble: 16 x 5

	carrier	mean_arr_delay	se_arr_delay	mean_dep_delay	se_dep_delay
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	9E	7.38	0.381	16.7	0.348
2	AA	0.364	0.238	8.59	0.209
3	AS	-9.93	1.37	5.80	1.18
4	B6	9.46	0.184	13.0	0.165
5	DL	1.64	0.203	9.26	0.182

6	EV	15.8	0.221	20.0	0.205
7	F9	21.9	2.36	20.2	2.23
8	FL	20.1	0.960	18.7	0.933
9	HA	-6.92	4.06	4.90	4.01
10	MQ	10.8	0.273	10.6	0.247
11	OO	11.9	9.02	12.6	8.00
12	UA	3.56	0.170	12.1	0.148
13	US	2.13	0.235	3.78	0.199
14	VX	1.76	0.699	12.9	0.626
15	WN	9.65	0.427	17.7	0.394
16	YV	15.6	2.27	19.0	2.11

```
# another approach
```

```
se <- function(x, na.rm = FALSE) {
  n <- length(x) - sum(is.na(x)) # makes the sample size accurate by removing the number of NA
  sqrt(var(x, na.rm = na.rm) / n)
}

se(c(1:5, NA), na.rm = TRUE)
```

```
[1] 0.7071068
```

```
flights |>
  group_by(carrier) |>
  summarize(
    mean_arr_delay = mean(arr_delay, na.rm = TRUE),
    se_arr_delay = se(arr_delay, na.rm = TRUE),
    mean_dep_delay = mean(dep_delay, na.rm = TRUE),
    se_dep_delay = se(dep_delay, na.rm = TRUE)
  )
```

```
# A tibble: 16 x 5
```

	carrier	mean_arr_delay	se_arr_delay	mean_dep_delay	se_dep_delay
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	9E	7.38	0.381	16.7	0.348
2	AA	0.364	0.238	8.59	0.209
3	AS	-9.93	1.37	5.80	1.18
4	B6	9.46	0.184	13.0	0.165
5	DL	1.64	0.203	9.26	0.182
6	EV	15.8	0.221	20.0	0.205

7 F9	21.9	2.36	20.2	2.23
8 FL	20.1	0.960	18.7	0.933
9 HA	-6.92	4.06	4.90	4.01
10 MQ	10.8	0.273	10.6	0.247
11 OO	11.9	9.02	12.6	8.00
12 UA	3.56	0.170	12.1	0.148
13 US	2.13	0.235	3.78	0.199
14 VX	1.76	0.699	12.9	0.626
15 WN	9.65	0.427	17.7	0.394
16 YV	15.6	2.27	19.0	2.11

6. Create `both_na()`, a function that takes two vectors of the same length and returns how many positions have an NA in both vectors. Hint: create two vectors like `test_x <- c(1, 2, 3, NA, NA)` and `test_y <- c(NA, 1, 2, 3, NA)` and write code that works for `test_x` and `test_y`, then turn it into a function that can handle any `x` and `y`. (In this case, the answer would be 1, since both vectors have NA in the 5th position.) Test it for at least one more combination of `x` and `y`.

```
both_na <- function(x, y) {
  sum(is.na(x) & is.na(y))
}

test_x <- c(1, 2, 3, NA, NA)
test_y <- c(NA, 1, 2, 3, NA)

test_a <- c(NA, NA, NA, 7, 21, 38)
test_b <- c(4, NA, NA, 21, NA, 78)

both_na(test_x, test_y)
```

```
[1] 1
```

```
both_na(test_a, test_b)
```

```
[1] 2
```

7. Run your code from (6) with the following two vectors: `test_x <- c(1, 2, 3, NA, NA, NA)` and `test_y <- c(NA, 1, 2, 3, NA)`. Did you get the output you wanted or expected? Modify your function using `if`, `else`, and `stop` to print an error if `x` and `y` are not the same length. Then test again with `test_x`, `test_y` and the sets of vectors you used in (6).

```
test_x2 <- c(1, 2, 3, NA, NA, NA)
test_y2 <- c(NA, 1, 2, 3, NA)

both_na(test_x2, test_y2)
```

Warning in is.na(x) & is.na(y): longer object length is not a multiple of shorter object length

```
[1] 2
```

```
both_na <- function(x, y) {
  if(length(x) != length(y)) {
    stop("`x` and `y` must be the same length", call. = FALSE)
  } else{sum(is.na(x) & is.na(y))
  }
}

both_na(test_x, test_y)
```

```
[1] 1
```

8. Here is a way to get not_cancelled flights in the flights dataset:

```
not_cancelled <- flights |>
  filter(!is.na(dep_delay), !is.na(arr_delay))
```

Is it necessary to check is.na for both departure and arrival? Using summarize, find the number of flights missing departure delay, arrival delay, and both. (Use your new function!)

It is not necessary to check for the not cancelled flights as there would be no missing values for those columns.

```
not_cancelled |>
  summarize(missing_dep = sum(is.na(dep_delay)),
            missing_arr = sum(is.na(arr_delay)),
            missing_both = both_na(dep_delay, arr_delay))
```

```
# A tibble: 1 x 3
  missing_dep missing_arr missing_both
    <int>         <int>         <int>
1         0           0           0
```

```
flights |>
  summarize(missing_dep = sum(is.na(dep_delay)),
            missing_arr = sum(is.na(arr_delay)),
            missing_both = both_na(dep_delay, arr_delay))
```

```
# A tibble: 1 x 3
  missing_dep missing_arr missing_both
    <int>         <int>         <int>
1     8255         9430         8255
```

9. Read the code for each of the following three functions, puzzle out what they do, and then brainstorm better names.

```
f1 <- function(time1, time2) {
  hour1 <- time1 %/% 100
  min1 <- time1 %% 100
  hour2 <- time2 %/% 100
  min2 <- time2 %% 100

  (hour2 - hour1)*60 + (min2 - min1)
}

f2 <- function(lengthcm, widthcm) {
  (lengthcm / 2.54) * (widthcm / 2.54)
}

f3 <- function(x) {
  fct_collapse(x, "non answer" = c("No answer", "Refused",
                                   "Don't know", "Not applicable"))
}
```

f1 overall finds the differences in times in minutes by calculating the remainder of the time (for both times) after being divided by 100 and labeling it hour, calculating the time divided by 100 and rounding down and labeling it min. Finally, the hours are subtracted from each other and multiplied by 60 and added to the difference in the minutes. A better name for the function would be `diff_time_min`.

f2 converts length and width in centimeters into inches and multiplies the two by each other to find the area of the object. A better name for this function would be `area_in`.

f3 collapses all responses that reveal no answer to the question into a “non answer” category. A better name for this function would be `collapse_na`.

10. Explain what the following function does and demonstrate by running `foo1(x)` with a few appropriately chosen vectors `x`. (Hint: set `x` and run the “guts” of the function piece by piece.)

```
foo1 <- function(x) {  
  diff <- x[-1] - x[1:(length(x) - 1)]  
  sum(diff < 0)  
}  
  
# figuring out the function  
x <- c(-11, 2, 33, 4, -55, 6, 77, 8, 99, 10)  
  
x[-1]
```

```
[1]  2 33  4 -55  6 77  8 99 10
```

```
x[1:length(x) - 1]
```

```
[1] -11  2 33  4 -55  6 77  8 99
```

```
diff
```

```
function (x, ...)  
UseMethod("diff")  
<bytecode: 0x1336423d0>  
<environment: namespace:base>
```

```
x <- c(-11, 2, 33, 4, -55, 6, 77, 8, 99, 10)  
foo1(x)
```

```
[1] 4
```

```
# vectors  
x <- c(-11, 2, 33, 4, -55, 6, 77, 8, 99, 10)  
foo1(x)
```

```
[1] 4
```

```
y <- c(-20, 17, 13, -19, -4, 7)
foo1(y)
```

```
[1] 2
```

foo1 first finds the difference between x with its first value removed and x with its last value removed. The function then finds the sum of differences that are less than 0.

11. The `foo1()` function doesn't perform well if a vector has missing values. Amend `foo1()` so that it produces a helpful error message and stops if there are any missing values in the input vector. Show that it works with appropriately chosen vectors `x`. Be sure you add `error = TRUE` to your R chunk, or else knitting will fail!

```
foo1 <- function(x) {
  if(is.na(sum(x))) {
    stop("`x` cannot have NAs", call. = FALSE)
  } else {
    diff <- x[-1] - x[1:(length(x) - 1)]
    sum(diff < 0)
  }
}

z <- c(1, NA, 2, 5, 17, NA)
foo1(z)
```

12. Write a function called `greet` using `if`, `else if`, and `else` to print out “good morning” if it's before 12 PM, “good afternoon” if it's between 12 PM and 5 PM, and “good evening” if it's after 5 PM. Your function should work if you input a time like: `greet(time = "2018-05-03 17:38:01 CDT")` or if you input the current time with `greet(time = Sys.time())`. [Hint: check out the `hour` function in the `lubridate` package]

```
library(lubridate)

greet <- function(time) {
  if(hour(time) < 12) {
    print("good morning")
  } else if(hour(time) >= 12 & hour(time) <= 17) {
    print("good afternoon")
  } else {
    print("good evening")
  }
}
```

```
}

greet(time = "2018-05-03 17:38:01 CDT")
```

```
[1] "good afternoon"
```

```
greet(time = "2018-05-03 18:38:01 CDT")
```

```
[1] "good evening"
```

```
greet(time = "2018-05-03 4:38:01 CDT")
```

```
[1] "good morning"
```

13. Modify the `summary6()` function from earlier to add an argument that gives the user an option to remove missing values, if any exist. Show that your function works for (a) the `hwy` variable in `mpg_tbl <- as_tibble(mpg)`, and (b) the `age` variable in `gss_cat`.

```
summary6 <- function(data, var, na.rm = FALSE) {
  data |> summarize(
    mean = mean({{ var }}, na.rm = na.rm),
    median = median({{ var }}, na.rm = na.rm),
    sd = sd({{ var }}, na.rm = na.rm),
    IQR = IQR({{ var }}, na.rm = na.rm),
    n = n(),
    n_miss = sum(is.na({{ var }})),
    .groups = "drop"      # to leave the data in an ungrouped state
  )
}

mpg_tbl <- as_tibble(mpg)
mpg_tbl |>
  summary6(hwy, na.rm = TRUE)
```

```
# A tibble: 1 x 6
   mean median    sd  IQR     n n_miss
<dbl> <dbl> <dbl> <dbl> <int> <int>
1  23.4     24  5.95     9   234     0
```

```
gss_cat |>
  summary6(age, na.rm = TRUE)
```

```
# A tibble: 1 x 6
  mean median    sd  IQR      n n_miss
<dbl> <int> <dbl> <dbl> <int> <int>
1  47.2     46  17.3   26 21483     76
```

14. Add an argument to (13) to produce summary statistics by group for a second variable (you should now have 4 possible inputs to your function). Show that your function works for (a) the `hwy` variable in `mpg_tbl <- as_tibble(mpg)` grouped by `drv`, and (b) the `age` variable in `gss_cat` grouped by `partyid`.

```
summary6 <- function(data, var, group, na.rm = FALSE) {
  data |>
    group_by({{ group }}) |>
    summarize(
      mean = mean({{ var }}, na.rm = na.rm),
      median = median({{ var }}, na.rm = na.rm),
      sd = sd({{ var }}, na.rm = na.rm),
      IQR = IQR({{ var }}, na.rm = na.rm),
      n = n(),
      n_miss = sum(is.na({{ var }})),
      .groups = "drop"      # to leave the data in an ungrouped state
    )
}

mpg_tbl |>
  summary6(hwy, drv, na.rm = TRUE)
```

```
# A tibble: 3 x 7
  drv    mean median    sd  IQR      n n_miss
<chr> <dbl> <dbl> <dbl> <dbl> <int> <int>
1 4     19.2     18  4.08     5  103     0
2 f     28.2     28  4.21     3  106     0
3 r     21      21  3.66     7   25     0
```

```
gss_cat |>
  summary6(age, partyid, na.rm = TRUE)
```


A tibble: 10 x 7

	partyid	mean	median	sd	IQR	n	n_miss
	<fct>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>
1	No answer	50.8	48	18.7	28	154	9
2	Don't know	34	34	NA	0	1	0
3	Other party	45.2	44.5	15.8	23	393	3
4	Strong republican	51.9	51	17.0	26	2314	8
5	Not str republican	47.2	45	17.2	26	3032	8
6	Ind,near rep	47.1	46	17.1	27	1791	2
7	Independent	43.3	41	16.3	24	4119	18
8	Ind,near dem	44.9	43	17.1	27	2499	2
9	Not str democrat	46.5	44	17.3	26.5	3690	11
10	Strong democrat	51.2	50	17.4	27	3490	15

15. Create a function that has a vector as the input and returns the last value. (Note: Be sure to use a name that does not write over an existing function!)

```
last_value <- function(x) {  
  x[length(x)]  
}  
  
x <- c("hello", "my", "name", "is", "gaby")  
last_value(x)
```

[1] "gaby"

```
numbers <- c(1:1111)  
last_value(numbers)
```

[1] 1111

16. Save your final table from (14) and write a function to draw a scatterplot of a measure of center (mean or median - user can choose) vs. a measure of spread (sd or IQR - user can choose), with points sized by sample size, to see if there is constant variance. Each point should be labeled with partyid, and the plot title should reflect the variables chosen by the user.

Hint: start with a ggplot with no user input, and then functionize:

- I don't know how to have the title be the variables selected without manually creating a title

```
library(ggrepel)

party_age |>
  ggplot(aes(x = mean, y = sd)) +
    geom_point(aes(size = n)) +
    geom_smooth(method = lm) +
    geom_label_repel(aes(label = partyid)) +
    labs(title = "Mean vs SD")
```

```
library(ggrepel)
party_age <-
  as_tibble (gss_cat |> summary6(age, partyid, na.rm = TRUE))

party_plot <- function(data, center, spread, size, label, title) {
  ggplot(data, aes(x = {{ center }}, y = {{ spread }})) +
    geom_point(aes(size = {{ size }})) +
    geom_smooth(method = lm) +
    geom_label_repel(aes(label = {{ label }})) +
    labs(title = title)
}

party_age |>
  party_plot(mean, sd, n, partyid, title = "Mean vs. SD")
```

`geom\_smooth()` using formula = 'y ~ x'

Warning: Removed 1 row containing non-finite outside the scale range
(`stat\_smooth()`).

Warning: Removed 1 row containing missing values or values outside the scale range
(`geom\_point()`).

Warning: Removed 1 row containing missing values or values outside the scale range
(`geom\_label\_repel()`).

